


AYOMI AI-Classes



AI (Artificial Intelligence) & ML (Machine Learning) Syllabus

1. <i>Introduction to Artificial Intelligence & Machine Learning</i>	<ul style="list-style-type: none"> • What is Artificial Intelligence? • Types of AI: Narrow AI vs. General AI • What is Machine Learning? • AI vs. Machine Learning vs. Deep Learning • Applications of AI and ML • Real-world examples of AI and ML
2. <i>Mathematics for Machine Learning and AI</i>	<ul style="list-style-type: none"> • Linear Algebra: <ul style="list-style-type: none"> ◦ Vectors, Matrices, ◦ Eigenvalues/Eigenvectors ◦ Matrix operations ◦ Singular Value Decomposition (SVD) • Probability & Statistics: <ul style="list-style-type: none"> ◦ Probability distributions (Normal, Binomial, Poisson, etc.) ◦ Bayes' Theorem ◦ Hypothesis Testing, p-values, confidence intervals • Calculus: <ul style="list-style-type: none"> ◦ Derivatives and Gradients ◦ Optimization techniques (Gradient Descent) • Optimization: <ul style="list-style-type: none"> ◦ Convex vs. Non-convex problems ◦ Stochastic Gradient Descent (SGD)
3. <i>Programming for AI and ML</i>	<ul style="list-style-type: none"> • Introduction to Python (or R) • Key Libraries: <ul style="list-style-type: none"> ◦ Python: NumPy, Pandas, Matplotlib, Seaborn, Scikit-learn, TensorFlow, Keras, PyTorch ◦ R: ggplot2, dplyr, caret • Data Structures and Algorithms for AI • Functions, Loops, Recursion, and Debugging • Software Development Practices for AI (Version control with Git)
4. <i>Supervised Learning Algorithms</i>	<ul style="list-style-type: none"> • Regression: <ul style="list-style-type: none"> ◦ Linear Regression

- Polynomial Regression
- Regularized Regression (Ridge, Lasso)
- Classification:
 - Logistic Regression
 - k-Nearest Neighbors (k-NN)
 - Support Vector Machines (SVM)
 - Decision Trees and Random Forests
 - Naive Bayes
 - Ensemble Methods (Bagging, Boosting, AdaBoost, Gradient Boosting)
- Model Evaluation:
 - Cross-validation
 - Accuracy, Precision, Recall, F1 Score
 - ROC-AUC curve
- Hyperparameter Tuning (Grid Search, Random Search)

5. Unsupervised Learning Algorithms

- Clustering:
 - k-Means Clustering
 - Hierarchical Clustering
 - DBSCAN
- Dimensionality Reduction:
 - Principal Component Analysis (PCA)
 - t-Distributed Stochastic Neighbor Embedding (t-SNE)
 - Autoencoders
- Anomaly Detection (Isolation Forest)

6. Deep Learning Audience

- Introduction to Neural Networks:
 - Perceptron and Feedforward Neural Networks
 - Activation Functions (ReLU, Sigmoid, Tanh)
 - Backpropagation and Gradient Descent
- Advanced Deep Learning Architectures:
 - Convolutional Neural Networks (CNNs) for Image Data
 - Recurrent Neural Networks (RNNs) for Sequential Data
 - Long Short-Term Memory (LSTM) Networks
- Transfer Learning and Fine-tuning Pre-trained Models
- Generative Models:

	<ul style="list-style-type: none"> ○ Generative Adversarial Networks (GANs) ○ Variational Autoencoders (VAE)
7. Natural Language Processing (NLP)	<ul style="list-style-type: none"> ● Text Preprocessing (tokenization, stemming, lemmatization) ● ● Bag-of-Words (BoW) Model ● Term Frequency-Inverse Document Frequency (TF-IDF) ● Word Embeddings (Word2Vec, GloVe, FastText) ● Sequence Models: <ul style="list-style-type: none"> ○ Recurrent Neural Networks (RNN) ○ LSTM, GRU ● Advanced NLP Models: <ul style="list-style-type: none"> ○ Transformer Networks (BERT, GPT, T5) ○ Attention Mechanisms and Self-Attention ● Applications of NLP: <ul style="list-style-type: none"> ○ Text Classification, Named Entity Recognition (NER), Sentiment Analysis
8. Reinforcement Learning	<ul style="list-style-type: none"> ● Introduction to Reinforcement Learning (RL) ● Key Concepts: Agent, Environment, Actions, Rewards ● Markov Decision Processes (MDPs) ● Q-Learning and Temporal Difference Learning ● Policy Gradient Methods ● Deep Reinforcement Learning (DRL) ● Applications of RL (Game Playing, Robotics)
9. Ethics and Bias in AI	<ul style="list-style-type: none"> ● AI Bias and Fairness ● Transparency and Explainability (Explainable AI) ● Ethical Considerations in AI ● Privacy and Security in Machine Learning
10. Advanced Topics (Optional)	<ul style="list-style-type: none"> ● Time Series Analysis and Forecasting <ul style="list-style-type: none"> ○ ARIMA Models, LSTM for time series prediction ● AI in Computer Vision <ul style="list-style-type: none"> ○ Object Detection and Segmentation (YOLO, SSD, Faster R-CNN)

	<ul style="list-style-type: none"> ○ Transfer Learning in Vision Models ● AI in Healthcare, Autonomous Vehicles, Robotics ● Federated Learning ● Meta-Learning ● Graph Neural Networks (GNNs)
<i>11. Model Deployment and Production</i>	<ul style="list-style-type: none"> ● Introduction to Model Deployment ● Model Deployment with Flask/Django ● Containerization with Docker ● Cloud Platforms (AWS, Google Cloud, Azure) ● Model Monitoring and Versioning
<i>12. Capstone Project</i>	<ul style="list-style-type: none"> ● Real-world problem solving using AI/ML ● End-to-End implementation of a model ● Documentation, Reporting, and Presentation